

Retrieval and De-duplication

Retrieval and De-duplication I

Retrieval and De-duplication II

Retrieval dispatches the configured query across `{{N_ENGINES}}` independent literature engines (`{{ENGINE_LIST}}`). Each engine is an isolated adapter exposing a uniform `search(query) -> list[Record]` interface; engines that are keyless — arXiv, OpenAlex [`@priem2022openalex`], Crossref [`@hendricks2020crossref`], PubMed/Entrez [`@sayers2022entrez`], SovietRxiv / RussiaRxiv, ChinaRxiv, Europe PMC, and bioRxiv/medRxiv — need no credentials, while Semantic Scholar [`@kinney2023semantic`] uses a key when present. SovietRxiv is a translated archive of Soviet-era scientific preprints sourced from Math-Net.Ru and CyberLeninka [`@sovietrxiv`]; ChinaRxiv serves translated Chinese preprints from ChinaXiv via the same unified API. Both retain original-language PDFs alongside each translation, and their polite rate-limit pool (300/min vs 30/min anonymous) is activated by an optional `X-API-Email` header. Europe PMC is a keyless biomedical aggregator covering PubMed, PMC, patents, and preprints in a single search call. bioRxiv/medRxiv share one unified date-window + cursor API; unlike the other engines it is not a free-text search endpoint so the adapter walks the date window page by page and

Engine Details I

Each engine adapter follows a uniform contract: a module-level API URL constant, a pure `_parse_*` parser function, and a `search_*` entry point with pagination, retry, and graceful error handling. All functions accept an injectable `base_url` parameter for hermetic testing with `pytest-httpserver` — no engine hardcodes its URL inside the function body.

Engine	Rate limit	Pagination	Auth
arXiv	3s between requests	100/page, offset	Keyless
Semantic Scholar	1 req/s (unauth.)	100/page, offset	Optional key
OpenAlex	Polite pool (mailto)	200/page, cursor	Keyless
Crossref	Polite pool (mailto)	1,000/page, offset	Keyless
PubMed	NCBI usage policy	restart/retmax	Keyless
SovietRxiv	30/min (300/min polite)	1–100/page, cursor	X-API-Email
ChinaRxiv	30/min (300/min polite)	1–100/page, cursor	X-API-Email
Europe PMC	~10 req/s (undocumented hard limit)	Up to 1,000/page	Keyless
bioRxiv/medRxiv	No documented limit	100/page fixed, cursor	Keyless

Engine Details II

Every new search writes `output/data/retrieval_report.json`, a timestamp-free report that records each attempted, skipped, or failed source with fetched, new-record, and duplicate counts. A zero-result response is therefore distinguishable from a disabled adapter or an HTTP failure. The committed corpus predates that report contract, so this paper intentionally does not reconstruct source-specific counts from the merged corpus.

Canonical Identifier Hierarchy I

Heterogeneous records are reconciled by a **canonical identifier hierarchy** — DOI > arXiv ID > Semantic Scholar ID > OpenAlex ID > a stable digest of the normalized title. When two records share a canonical identifier they are merged, keeping the version with the most complete metadata (a count of non-None optional fields). The DOI is normalized: case-folded, resolver-prefix stripped, so the same paper returned by two engines under case/format-variant DOIs merges. For this run, $\{\{\text{DOI_COUNT}\}\}$ records carry DOIs, $\{\{\text{OPENALEX_ID_COUNT}\}\}$ carry OpenAlex IDs, and $\{\{\text{ARXIV_ID_COUNT}\}\}$ carry arXiv IDs. The de-duplicated corpus for this run holds $N = \text{CORPUS_SIZE}$ records published across $\{\{\text{YEAR_START}\}\} - \{\{\text{YEAR_END}\}\}$.

Relevance Filtering I

After de-duplication, a relevance filter drops papers whose title and abstract contain none of the configured relevance keywords (`{{KEYWORDS_RELEVANCE}}`). Keywords are matched case-insensitively; an empty keyword list is treated as no filter to avoid silently wiping the corpus. A year filter then excludes papers published before the configured start year (`{{YEAR_START}}`).